

# *Asynchronous Consensus-Free Transaction Systems*



*Shoma Mori, Roland Schmid, Jakub Sliwinski, Prof. Roger Wattenhofer*

# Motivation

Q. Why need a consensus?

A. The systems must reach only one result


However

Solving a consensus costs much (e.g. PoW)

→ Consensus-free systems are desired

With PoW, the history of transactions can be overturned

→ Deterministic systems are also desired








Purpose

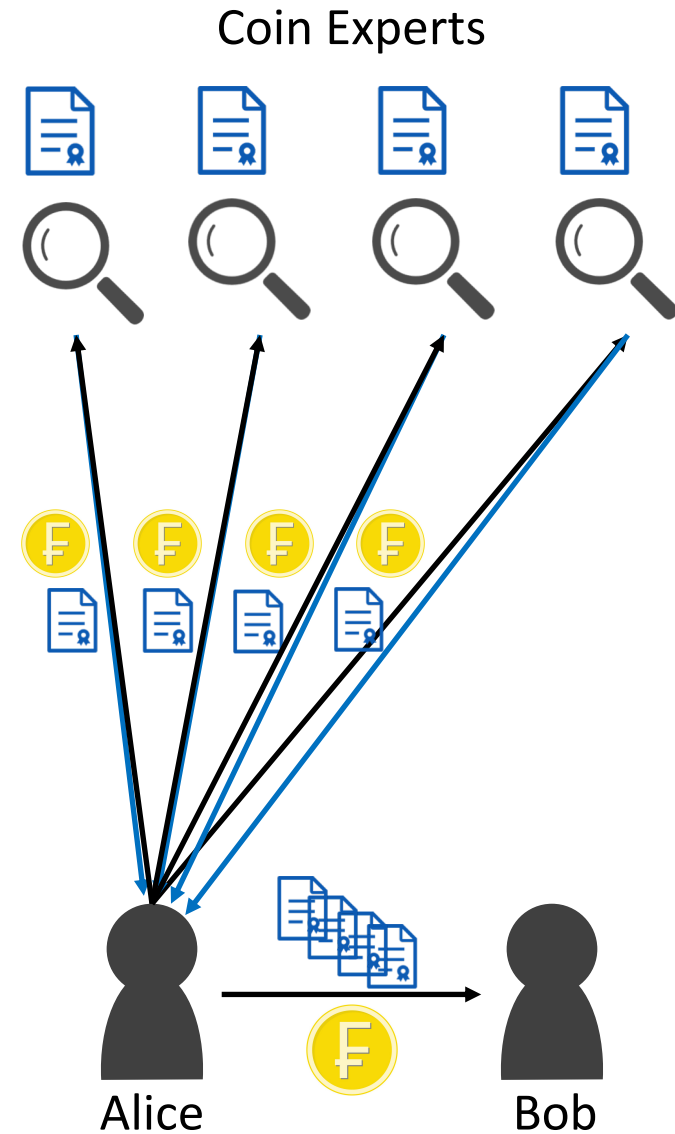
Building high performance systems without a consensus

- transaction/sec
- scalability

# Simple Overview of ACFTS

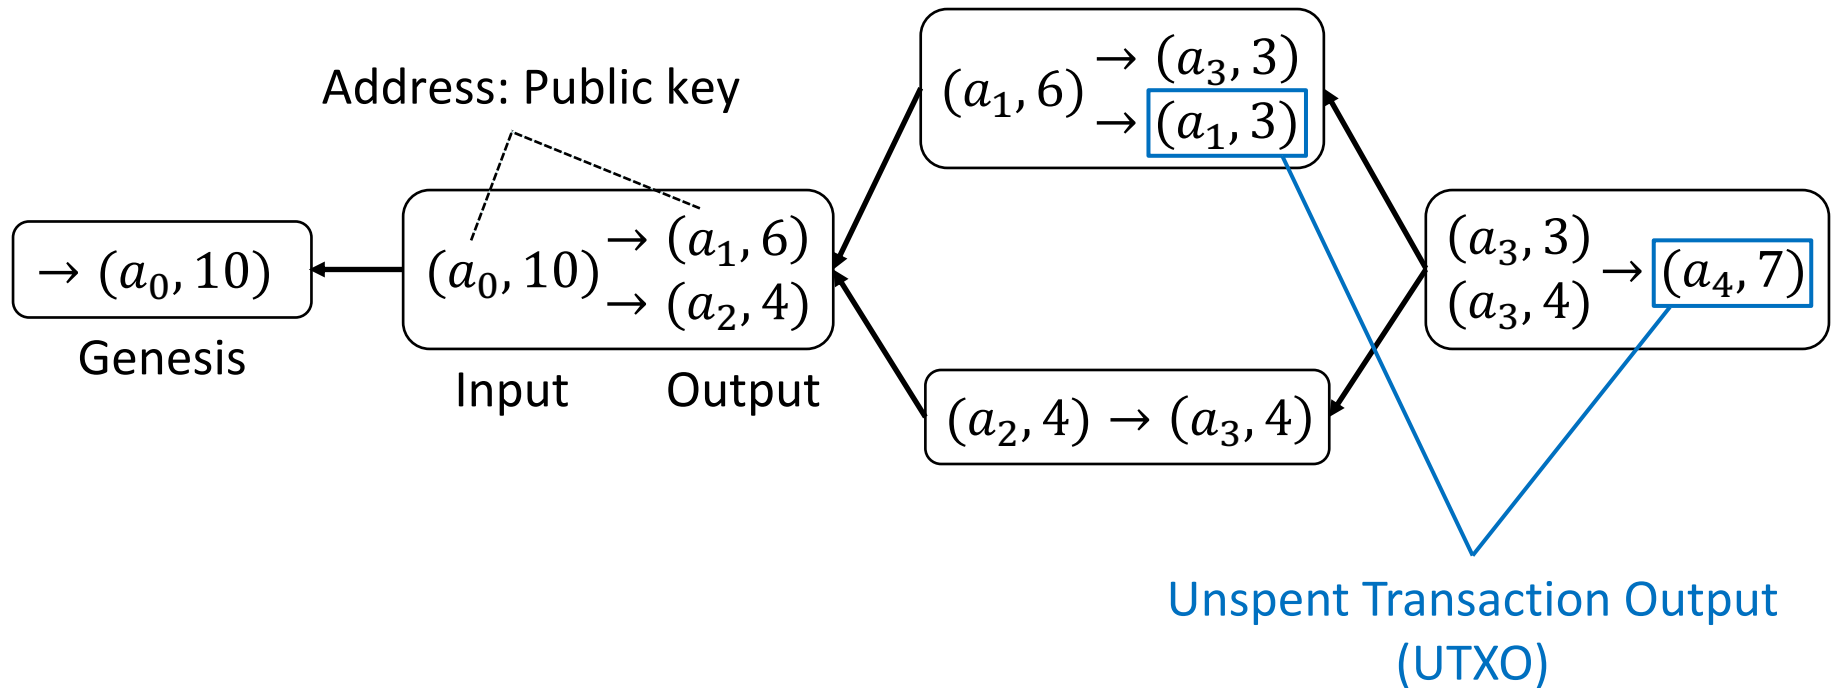
Example: Alice  $\xrightarrow{1\text{CHF}}$  Bob

- ① Alice searches 1CHF coin for her wallet
- ② Alice sends  to coin experts
- ③ The experts identify if  is real
- ④ The experts send back  to Alice
- ⑤ Alice sends Bob  with 



# Transaction in ACFTS

Transaction = Transfer of cryptocurrency



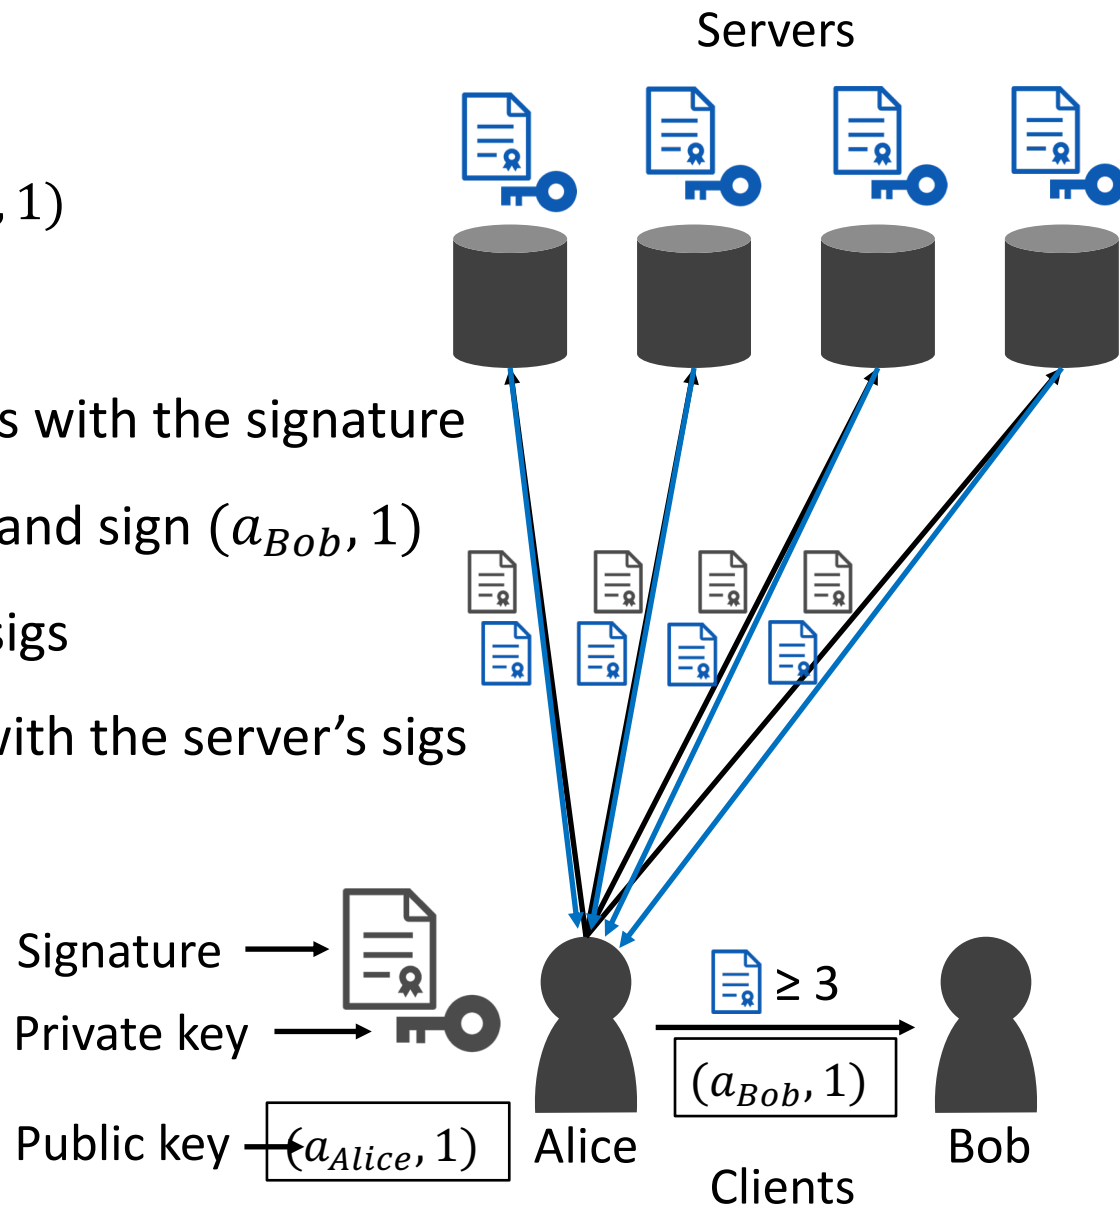
How does ACFTS avoid conflicting transactions?

→ "Servers" manage spending of transactions

# Overview of ACFTS

Example:  $(a_{Alice}, 1) \rightarrow (a_{Bob}, 1)$


- ① Alice signs  $(a_{Alice}, 1)$
- ② Alice sends servers requests with the signature
- ③ The servers verify the reqs and sign  $(a_{Bob}, 1)$
- ④ The servers send back the sigs
- ⑤ Alice sends Bob  $(a_{Bob}, 1)$  with the server's sigs




The number of  must be more than **2/3** of all servers<sup>[1]</sup>

# Client


Example case:




Alice's storage



Create  
signatures




More than 2/3 of all server's signatures  
to use each UTXO



Request

Change Output

# Server's verification



Request

## Verification process

- ✓ Does each input have the owner's sigs?
- ✓ Does each input have enough server's sigs?
- ✓ Are the UTXOs unused?

The status of all outputs

Output	Used
$(a_{Bob}, 2)$	false
$(a_{Bob}, 3)$	false
$(a_{Bob}, 4)$	false
$(a_{Bob}, 2)$	false
$\vdots$	$\vdots$

# Server's signing

If the request is valid,

- sign a hash of **all outputs**

$\rightarrow (a_{Alice}, 10)$   
 $\rightarrow (a_{Bob}, 1)$

**Not each output**

$\rightarrow$  For efficiency



Therefore

to use  $(a_{Alice}, 10)$ ,

Alice needs to send  and  $(a_{Bob}, 1)$

- update the status table
- send the signature to the client


The status of all outputs

Output	Used
$(a_{Bob}, 2)$	<del>true</del>
$(a_{Bob}, 3)$	<del>true</del>
$(a_{Bob}, 4)$	<del>true</del>
$(a_{Bob}, 2)$	<del>true</del>
$(a_{Alice}, 10)$	false
$(a_{Bob}, 1)$	false
$\vdots$	$\vdots$




# Implementation

- Communication protocol: **HTTP**



- Clients and servers manage **relational databases**
- Cluster
  - One cluster can own multiple addresses




Demo

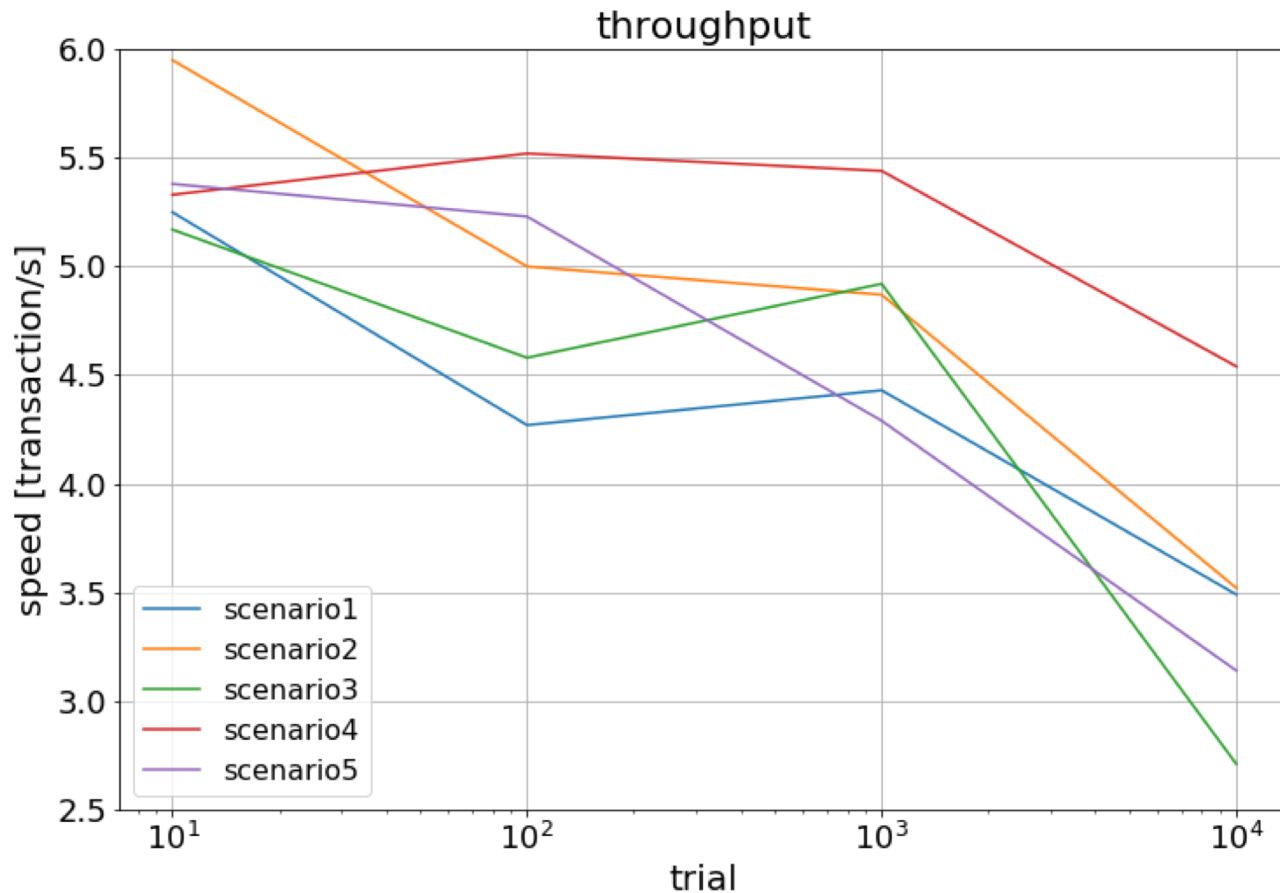
# Experiment

- Purposes
  - Measuring the throughputs (transaction/sec)
  - Finding bottlenecks
- Environment
  - 4 servers
  - 2 clients (clusters)
- Scenarios [“→”: 1 amount transfer, genesis:  $(a_A^0, 100000)$ ]

Scenario: 1



# Results







- Approximately 2.7–6.0 transaction/sec
  - Generally, the throughputs go down as increasing trial
- ➡ What are the bottlenecks preventing scalability?

# Bottlenecks

Profile [scenario1, trials = 10000]

## CPU usage in the client-side

name	first 60s [%]		last 60s [%]
getClientSig	43.4		40.4
updateOutputs	41.0		39.9
getServerSigs	6.13		10.3
findUTXOs	3.06		4.23

→ Sign a UTXO




→ Add new outputs & server's sigs to DB

→ **Search the DB** for server's sigs

→ **Search the DB** for available UTXOs

Client: Searching the DB does not seem to scale well

## CPU usage in the server-side

name	first 60s [%]		last 60s [%]
verifyUTXO	65.1		64.1
unlockUTXO	22.7		23.1
createSignature	10.9		10.5

→ Verify server's sigs on inputs

→ Verify client's sigs on inputs

→ Sign an output

Server: All functions seem to scale

Both: **Cryptographic processes** (verification & signing) are bottlenecks

# Discussion

Use a hash table for the server's signatures for scalability

Relational database

id	output	output_id	signature
1	$(a_{Alice}, 2)$	1	3eqr42
2	$(a_{Alice}, 4)$	1	a7eq6i
3	$(a_{Alice}, 3)$	1	3qu8iu
		2	1qer5i
		2	4qit6u

$O(n)$



Hash Table

key	data
$(a_{Alice}, 2)$	[3eqr42, a7eq6i, 3qu8iu]
$(a_{Alice}, 4)$	[1qer5i, 4qit6u]
$(a_{Alice}, 3)$	[]

$O(1)$

# Conclusion

- ACFTS can avoid conflicting transactions without a consensus
- Verification & signing are bottlenecks
- In the client-side, searching the DB does not scale well

## Future Work

- Change how to store outputs for scalability
- Reduce time for cryptographic processes
- Run in a real environment
- Adapt to the replacement of servers



# Hash table for findUTXOs

A hash table is not that effective for finding available UTXOs

Relational database

Need to store  
sibling outputs

id	output
1	$(a_{Alice}, 2)$
2	$(a_{Bob}, 1)$
3	$(a_{Alice}, 3)$
4	$(a_{Bob}, 2)$
5	$(a_{Alice}, 4)$




Hash Table


key	data
$a_{Alice}$	$(a_{Alice}, 2), (a_{Alice}, 3), (a_{Alice}, 4)$
$a_{Bob}$	$(a_{Bob}, 1), (a_{Bob}, 2)$

$O(m)$

Why more than  $2/3$ ?



Shuffle



$N - f$  nodes are required for quorum

